

COMP7940

Cloud Computing

Chapter 01

Characterization of Distributed Systems

Learning Outcomes


- Understand the concept of distributed systems
- Explain why we need distributed systems
- Give examples of distributed systems
- Understand the challenges in designing distributed systems

Outline

- Introduction
- Definition of Distributed Systems
- Example Distributed Systems
- Challenges of Distributed Systems
- Summary

Introduction

- Networks of computers are everywhere!
 - Global Internet
 - Mobile phone networks
 - Corporate networks
 - Factory networks
 - Campus networks
 - Home networks
 - In-car networks
 - On board networks in planes and trains
- Benefits of networks of computers:
 - Resource sharing
 - Hardware resources
 - File, database, etc.
 - Information: text, image, audio, video



Defining Distributed Systems (DS)


- "A system in which hardware or software components located at networked computers communicate and coordinate their actions only by message passing." -- from REF1
- Examples of Distributed Systems:
 - Cluster:
 - A type of parallel or distributed processing system, which consists of a collection of interconnected stand-alone computers cooperatively working together as a single, integrated computing resource
 - Science Faculty's cluster:
 - <http://site.sci.hkbu.edu.hk/hpccc/>
 - Cloud:
 - A type of parallel and distributed system consisting of a collection of interconnected and virtualised computers that are dynamically provisioned and presented as one or more unified computing resources based on service-level agreements established through negotiation between the service provider and consumers
 - <http://labs.hol.vmware.com/HOL/catalogs/catalog/681>

Motivation of Distributed Systems

- To share resources and information. E.g.,
 - Staff in an office share a printer through a LAN.
 - A group of students share files through a NAS server.
 - A company shares its information with the public through its Web server and Internet.
 - Tens of millions of users share storage resources at Google through Google Drive.
 - Billions of users share personal information through Facebook.

Reasons for Distributed Systems

- Functional Separation:
 - Existence of computers with different capabilities and purposes:
 - Clients and Servers
 - Data collection and data processing
- Inherent distribution:
 - Information:
 - Different information is created and maintained by different people (e.g., Web pages)
 - People
 - Computer supported collaborative work (virtual teams, engineering, virtual surgery)
 - Retail store and inventory systems for supermarket chains
- Power imbalance and load variation:
 - Distribute computational load among different computers.
- Reliability:
 - Long term preservation and data backup (replication) at different locations.
- Economies:
 - Sharing a printer by many users and reduce the cost of ownership.
 - Building a supercomputer out of a network of computers.




Examples of Distributed Systems

- They (DS) are based on familiar and widely used computer networks:

- Internet
- Intranets, and
- Wireless networks

- Example DS:

- Web (and many of its applications like Facebook)
- Data Centers and Clouds
- Wide area storage systems
- Banking Systems



Example 1: Web Search

- Web search engines such as google.com and baidu.com are critical to our work, study, and life.
- Web search is a very complex task
 - To index the entire contents of the World Wide Web, including web pages, multimedia sources, scanned books, etc.
 - Over 63 billion pages and one trillion unique web addresses (2011)
- A complicated distributed system is designed to provide web search service
 - A very large number of networked computers located at many data centres
 - A distributed file system to support very large data storage and processing
 - An associated structured distributed storage system for fast access to very large datasets
 - A lock service that offers distributed system functions such as distributed locking and agreement
 - A programming model for managing very large parallel and distributed computations

Example 2: Massively Multiplayer Online Games (MMOGs)

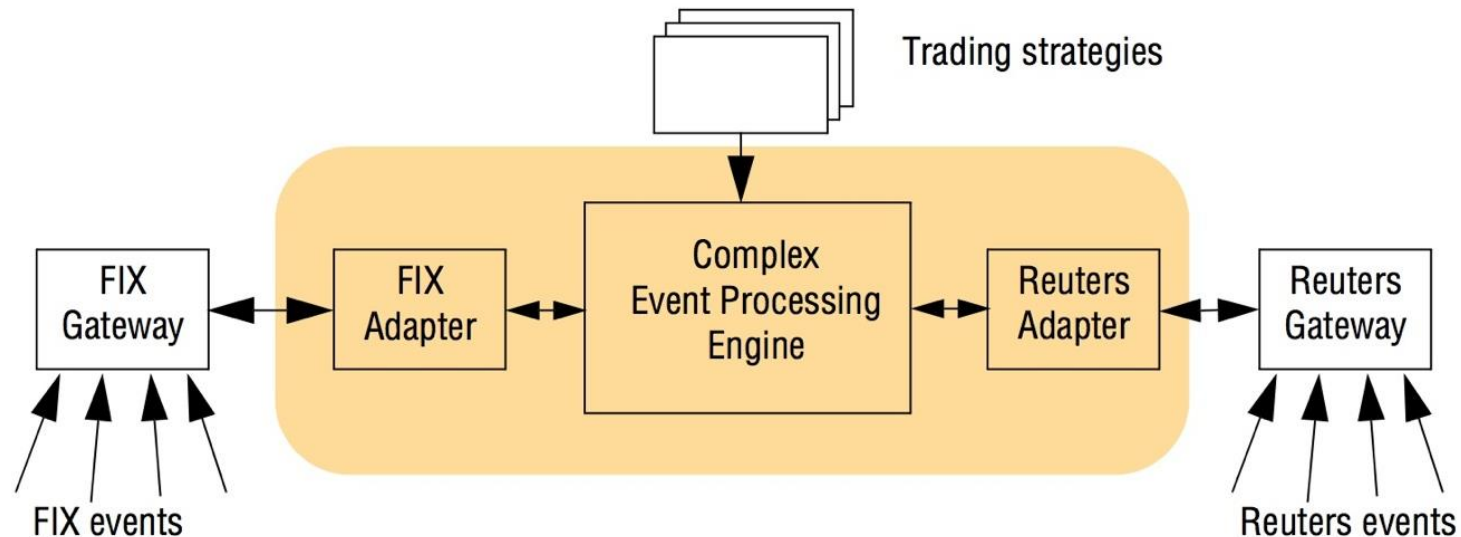
- MMOGs have the following technical challenges
 - Fast response time for good user experience
 - Real-time propagation of events to many players
 - How to maintain a consistent view of the shared world?
- Possible solutions:
 - Client-server architecture: a single copy of the state of the game world is maintained by a cluster of computers
 - The load is partitioned by allocating individual “star systems” to particular computers within the cluster.
 - Ensure fast user response through optimizing network protocols and rapid response to incoming events.
 - Distributed architecture: the universe is partitioned across many servers which may also be geographically distributed
 - Users are dynamically allocated a particular server based on current usage patterns and network delays to the server.

Example 3: Financial Trading

- Financial industry needs cutting edge distributed systems for real-time access to a wide range of information such as share prices and trends, and economic & political events.
- Distributed event-based systems are used to make intelligent trading strategies
 - Many suppliers provide different types of financial events (share price movements, interest rate changes, unemployment rate changes, etc.)
 - Many users subscribe to the suppliers to receive the events reliably in a timely manner

Example 3 (Cont.)

REF1 Figure 1.2. An example of financial trading system



- Heterogeneity: event sources are in a variety of formats
 - Example events: Reuter market data events, FIX events, etc.
 - FIX: Financial Information eXchange protocol
- Adapters are used to translate heterogeneous formats into a common internal format.
- The event processing engine needs to handle the incoming event streams in real-time, e.g., looking for patterns that indicate a trading opportunity (e.g., HSBC's stock prices in HK and London are slightly different due to exchange rate).

Example 3 (Cont.)

- Sample trading script (from www.progress.com):

WHEN

MSFT price moves outside 2% of MSFT Moving Average

FOLLOWED-BY (

MyBasket moves up by 0.5%

AND

HPQ's price moves up by 5%

OR

MSFT's price moves down by 2%

)

ALL WITHIN

any 2 minute time period

THEN

BUY MSFT

SELL HPQ

MSFT: Microsoft

HPQ: Hewlett Packard

Consequences of Distributed Systems

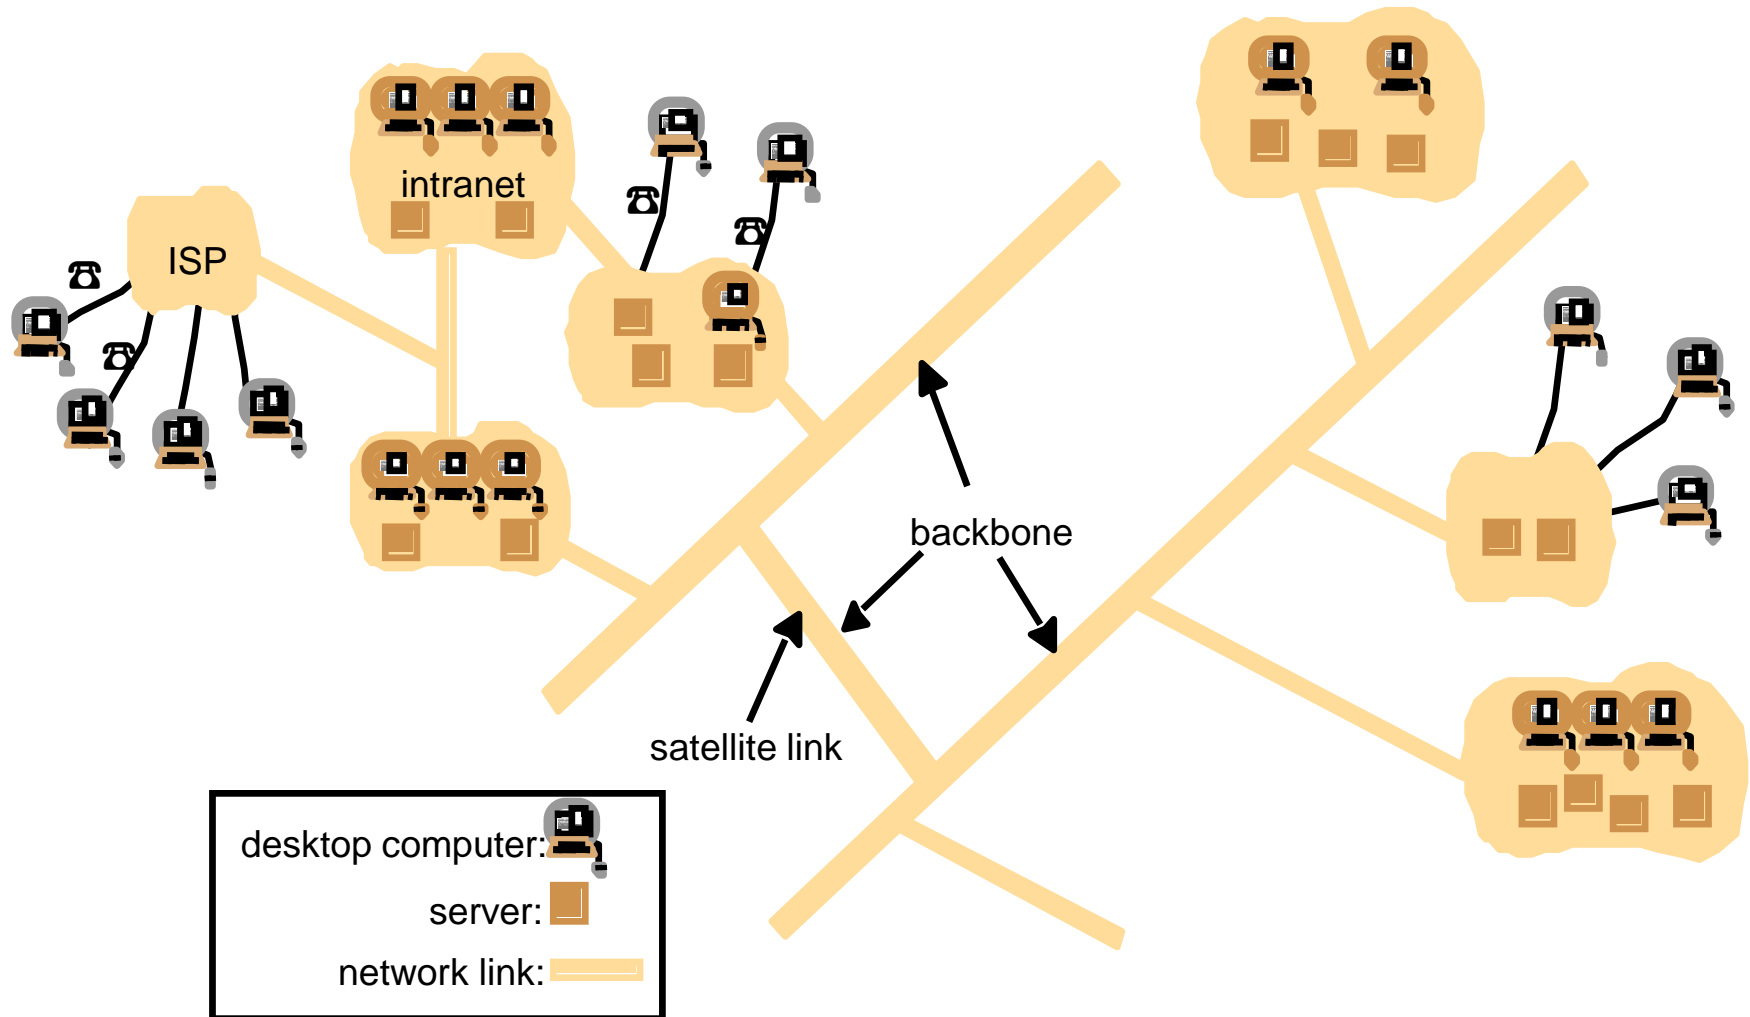
- Computers in distributed systems may be on separate continents, in the same building, or the same room. DSs have the following consequences:
 - Concurrency – each computer is autonomous.
 - Carry out tasks independently
 - Tasks coordinate their actions by exchanging messages.
 - System capacity can be increased by adding more resources.
 - No global clock
 - The clocks in different computers are different. The accuracy of clock synchronization is limited.
 - Independent Failures
 - Each component (network, computer, etc.) of a distributed system can fail independently,
 - but the other components do not know about the crash.

Trends in Distributed Systems

- Distributed systems are undergoing significant change with the following influential trends
 - The emergence of pervasive networking technology
 - The emergence of mobile and ubiquitous computing
 - The increasing demand for multimedia services
 - The view of distributed systems as a utility

Pervasive Networking and the modern Internet

REF Fig. 1.3 A typical portion of the Internet



Pervasive Networking and the modern Internet (cont.)

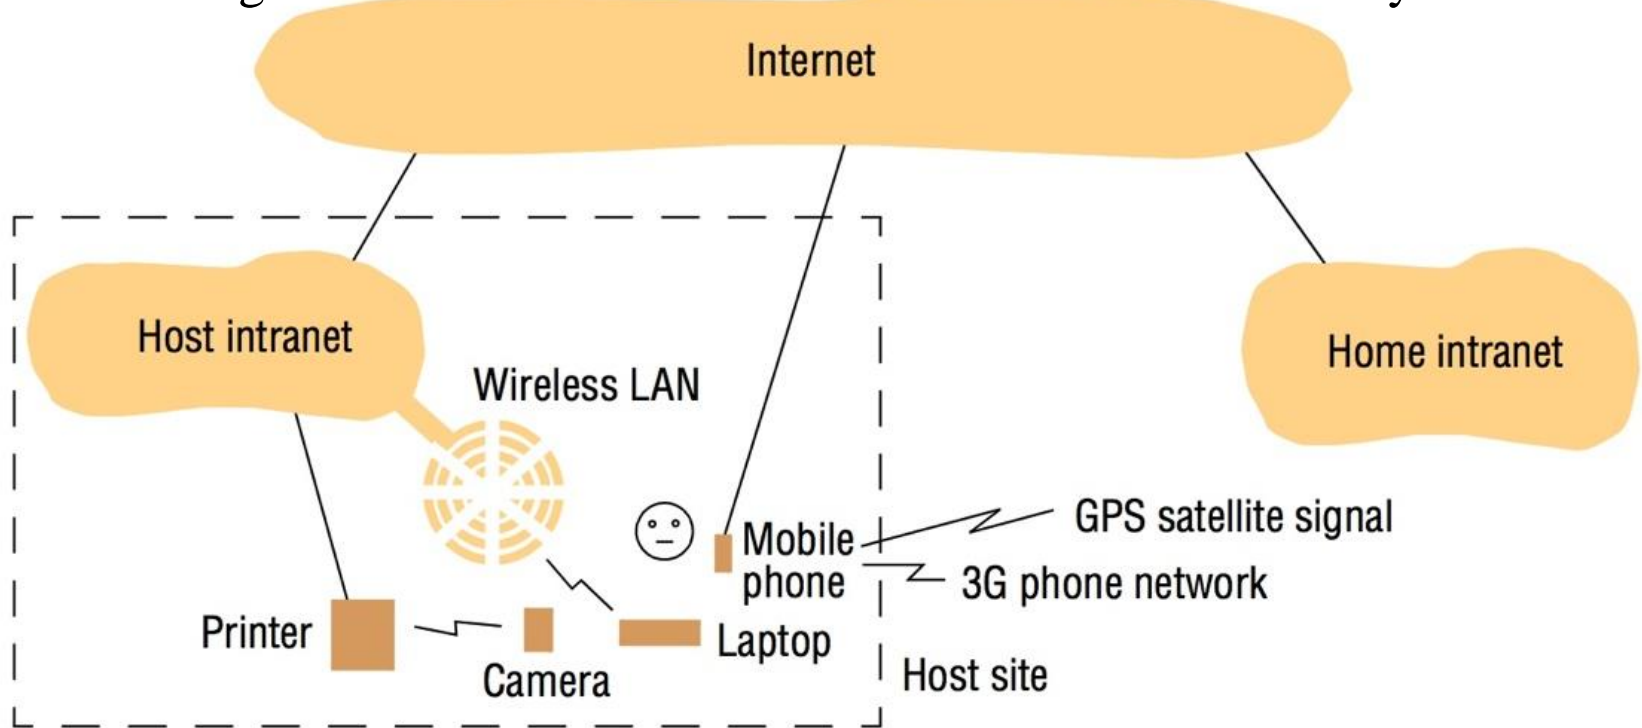
- Pervasive: spread throughout
- The Internet is a vast collection of computer networks of many different types and hosts various types of services.
 - Wireless: WiFi, WiMAX, Bluetooth, 4G or 5G, satellite, etc.
 - Wired: electronic, optical, tradition phone network, etc.
- Maintenance of intranet often an issue in distributed systems
 - No risk if no connection to internet, e.g., some intranets of police or security or law enforcement agencies are not connected to the internet.
 - Firewalls are used to limit services from/to an Intranet, e.g., no FTP or Remote Desktop access is allowed to intranet.
 - Firewalls are problematic by impeding legitimate access to services when resource sharing between internal and external users is required.

Mobile and Ubiquitous Computing

- Mobile computing: performing computing tasks while the user is on the move, away from his/her usual environment
 - Mobile devices: mobile phones, laptops, cameras, watches, etc.
 - Mobile devices can still access resources in home intranet (e.g., looking up a photo in your home hard disk via HKBU WiFi).

Mobile and Ubiquitous Computing (cont.)

REF Fig. 1.4 Portable and handheld devices in a distributed system



- A mobile phone user (e.g., student) using a host intranet (e.g., on HKBU campus).
- He can still connect to his home intranet.

Mobile and Ubiquitous Computing (cont.)

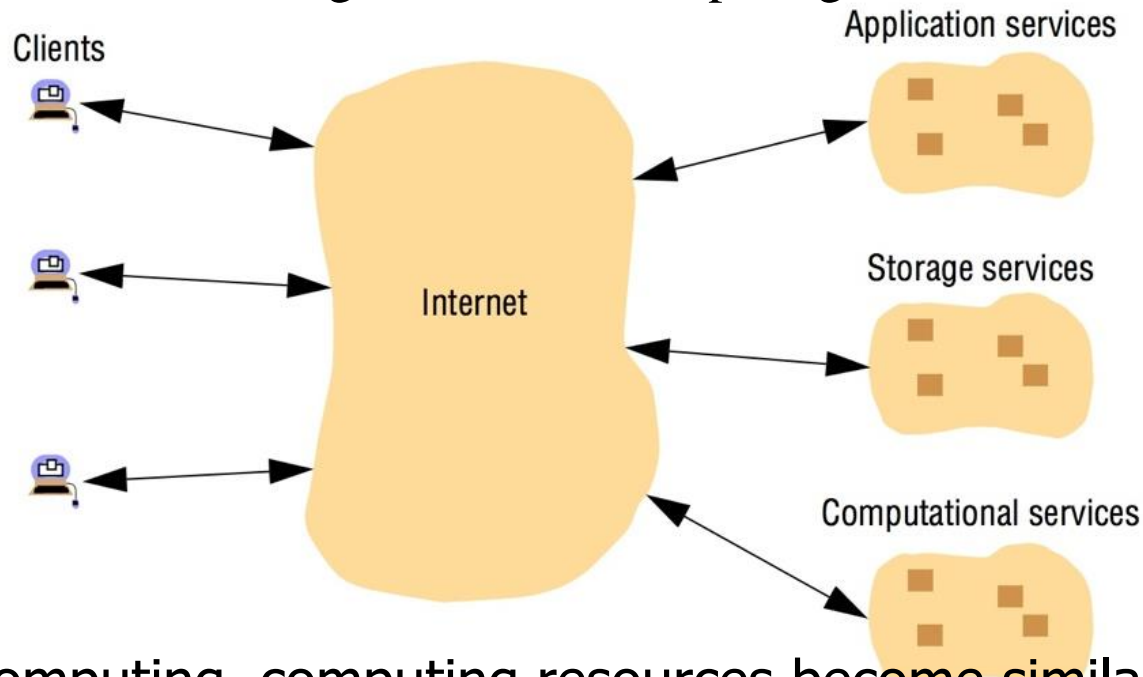
- Ubiquitous computing is the harnessing of many small, cheap computational devices in users' physical environments, including home, office and natural settings.
 - Printers, TVs, washing machines, Hi-Fi systems, cars, refrigerators, door locks, etc.
 - Needs to access resources conveniently located nearby (location-aware or context-aware computing).
- Ubiquitous and mobile computing may overlap.
 - E.g., At HKBU, you use your mobile phone (mobile computing) to print a pdf file in your phone using a lab printer (ubiquitous computing).
- Needs to support spontaneous interoperation.
 - Associations between devices are routinely created and destroyed.
 - Automatic service discovery

Distributed Multimedia Systems

- Needs to support the storage, transmission, and presentation of discrete media types, e.g., text, pictures, video, etc.
 - For each type, support different encoding and decoding formats (e.g., .mov, .wmv, .avi etc.)
- To broadcast across a group of users.
- Must preserve real-time relationship between elements of a media type.
 - Preserve frame sequence
 - Maintain minimum speed (# of frames per sec)
 - Maximum delay or latency
 - Quality of service (QoS)

Distributed Computing as a Utility

REF Fig. 1.5 Cloud Computing



- In cloud computing, computing resources become similar to other utilities such as water or electricity.
 - Resources: hard disk space, RAM, CPU, virtual network, virtual machine, virtual OS (key enabling technology of cloud), virtual firewall, software such as gmail, Google Map/Earth, etc.
- Cloud service providers: Amazon, Google, Alibaba, Microsoft, IBM, etc.
 - Big business

Challenges of Distributed Systems: An Example

- Suppose you own a web-based online store (such as amazon.com)
 - Customers can connect their computer to the web server:
 - Browse product catalog
 - Search and compare
 - Place orders
 - ...

List of Challenges

- Heterogeneity
 - Heterogeneous components must be able to interoperate
- Openness
 - Interfaces should be publicly available to ease inclusion of new components
- Security
 - The system should only be used in the way intended
- Scalability
 - System should work efficiently with an increasing number of users
 - System performance should increase with inclusion of additional resources

List of Challenges (Cont.)

- Fault handling
 - Detecting failures/Masking failures/Tolerating failures/Recovery from failures/Redundancy
- Concurrency
 - Shared access to resources must be possible
- Distribution transparency
 - Distribution should be hidden from the user as much as possible

Challenges I

- What if
 - Your company connects to the internet using optical fibers, but a customer connects to your website from an old-fashioned telephone network in a 3rd world country?
 - Your customer uses a completely different hardware?
 - PC, MAC, mobile phone...
 - Mobile code: program code that can be transferred from one computer to another and run at the destination, e.g., Java applets running on Java virtual machines.

Challenges I (con't)

- Your customer uses a completely different operating system?
 - Windows, Unix, Android, IOS...
 - All have implementations of the Internet protocols, but
 - interfaces (function calls) for exchanging messages are different in different OS.
- Your customer uses a different way of representing data?
 - ASCII, EBCDIC, ...
 - big endian, little endian integers
- **Heterogeneity**

Challenges II

- When building the system...
 - Do you want to write the whole software on your own (network, database,...)?
 - What about updates, new technologies?
 - **Reuse and Openness** (Standards)
 - The openness of distributed systems is determined by the degree to which new resource-sharing services can be added and made available for use by a variety of client programs.
 - Cannot be achieved unless the specification and documentation of the key software interfaces of the components are made available to software developers (published).
 - E.g., file transfer, email, telnet, etc.
 - Open distributed system can be constructed from heterogeneous hardware and software. The conformance of each component to the published standard must be carefully tested.

Challenges III


- What if
 - Someone tries to break into your system to steal data?
 - ... sniffs for information?
 - ... your customer orders something and does not accept the delivery saying he did not?
 - Security**
 - Confidentiality, integrity, authenticity, availability

Security I

- Resources are accessible to authorized users and used in the way they are intended
- Confidentiality
 - Protection against disclosure to unauthorized individual information
 - E.g. ACLs (access control lists) to provide authorized access to information
- Integrity
 - Protection against alteration or corruption
 - E.g. changing the account number or amount value in a money order

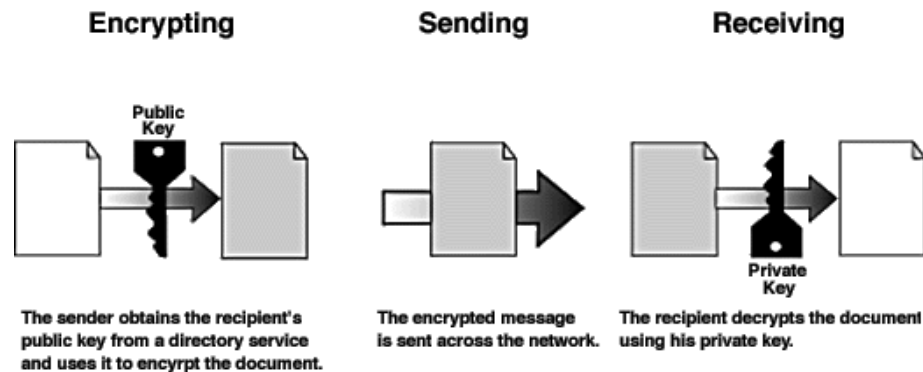
Security II

- Availability
 - Protection against interference targeting access to the resources.
 - E.g. denial of service (DoS, DDoS) attacks
- Authenticity or Non-repudiation
 - Proof of sending / receiving an information
 - E.g. digital signature



Security Mechanisms

- Encryption
 - E.g. Blowfish, AES, RSA
- Authentication
 - E.g. password, biometrics, public key authentication
- Authorization
 - E.g. access control lists




Challenges IV

- What if
 - You are so successful that millions of people are visiting your online store at the same time?
 - **Scalability**
 - A system is scalable if it remains effective when there is a significant increase in the number of resources and the number of users.
 - Controlling the cost of physical resources
 - For a system with n users to be scalable, the quantity of physical resources required should be at most $O(n)$.
 - Controlling the performance loss
 - Consider the management of a set of data whose size is proportional to the number of users or resources.

Big O notation

Usually Big O is used to measure time-complexity, but is also possible to measure other resource



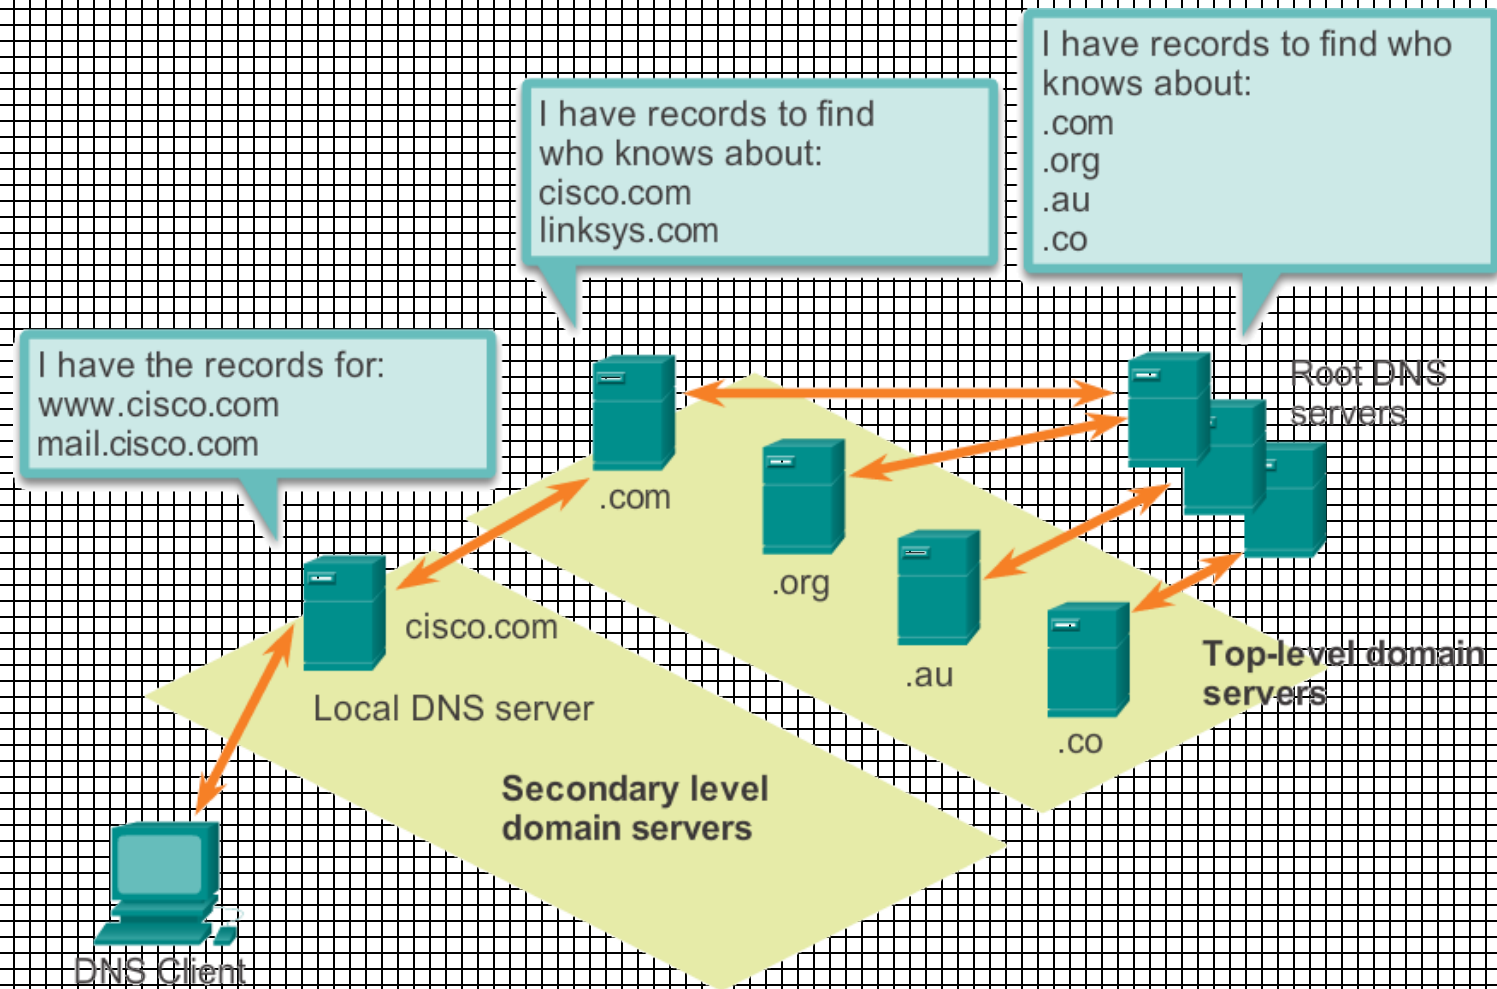
<https://dev.to/b0nbon1/understanding-big-o-notation-with-javascript-25mc>

Challenges IV

—Scalability (cont.)

- Preventing software resources running out
 - E.g., IP addresses (IPv4) using 32 bits are running out.
 - Now migrating to IPv6 using 128 bits – not easy task, a lot of software code must be changed.
- Avoiding performance bottlenecks
 - Algorithms should be decentralized.
 - Example
 - The predecessor of the DNS keeps the name table in a single master file that could be downloaded to any computer – a performance bottleneck.
 - The DNS removed this bottleneck by partitioning the name table between servers located throughout the Internet.

DNS Server




Challenges V

- What if
 - Two customers want to order the same item at the same time, or to bid on the same item in eBay?
 - **Concurrency**
 - avoid
 - Racing Condition
 - deadlock

Thread A		Thread B		Count
Instruction	Register	Instruction	Register	
LOAD Count	10	LOAD Count	10	10
ADD #1	11			10
STORE Count	11			11
		SUB #1	9	11
		STORE Count	9	9

<https://pages.mtu.edu/~shene/NSF-3/e-Book/RACE/overview.html>



<https://robertostefanettinavblog.com/2018/04/11/detect-deadlocks-in-old-navs-2017/>

Concurrency

- Provide and manage concurrent access to shared resources:
 - Fair scheduling
 - Preserve dependencies (e.g. distributed transactions)
 - Avoid deadlocks
 - Object locking, data consistency, semaphores

Java Concurrency




Challenges VI

- What if
 - The server or database with your inventory information crashes?
 - Your customer's computer crashes in the middle of an order?
 - The order is lost or corrupted during transmission?
 - **Fault tolerance**

Fault Tolerance

- Failure: an offered service no longer complies with its specification
- Fault: cause of a failure (e.g. crash of a component)
- Fault tolerance: no failure despite faults



Fault Tolerance Mechanisms

- Fault detection
 - Checksums, heartbeat (originator sends messages periodically), ...
- Fault masking
 - Retransmission of corrupted messages, redundancy, ...
 - Redundancy in
 - hardware, routing path, Domain Name System, databases, etc.
- Fault toleration
 - Exception handling, timeouts, ...
- Fault recovery
 - Rollback mechanisms, ...

Challenges VII

- What if
 - You want to move your business and servers to a different city (because of the weather)?
 - **Distribution or mobility transparency**
 - Transparency is defined as the concealment from the user and application programmer of the separation of components in a distributed system.
- Eight forms of transparency (as identified by Advanced Networked Systems Architecture ANSA, an industry standard)
 - Access transparency enables local and remote resources to be accessed using identical operations.
 - Location transparency enables resources to be accessed without knowledge of their physical or network location.

Challenges VI

- Eight forms of transparency (cont.)
 - Concurrency transparency enables several processes to operate concurrently using shared resources without interference between them.
 - Replication transparency enables multiple instances of resources to be used to increase reliability and performance without knowledge of the replicas by users or application programmers.
 - Failure transparency enables the concealment of faults.
 - Mobility transparency allows the movement of resources and clients within a system without affecting the operation of users or programs.
 - Performance transparency allows the system to be reconfigured to improve performance as loads vary.
 - Scaling transparency allows the system and applications to expand in scale without change to the system structure or the application algorithms.

Summary

- Distributed Systems are everywhere
- The Internet enables users throughout the world to access its services wherever they are located
- Resource sharing is the main motivating factor for constructing distributed systems
- Construction of DS produces many challenges:
 - Heterogeneity, Openness, Security, Scalability, Failure handling, Concurrency, and Transparency

Sample questions

- Give examples of heterogeneous computing devices in DS.
- Give a characteristic of open system.
- Give an example of security requirement in DS.
- Why scalability has been a challenge to DS?
- Explain a mechanism to handle failures in DS.
- What are challenges in concurrent computation in DS.
- Explain X transparency.